

## CS150A Quiz 2 Solutions

Assume that each page in our system can hold 64 KB (1 KB = 1024 bytes), integers are 32-bits wide, and bytes are 8-bits wide.

Consider the following relation:

```
CREATE TABLE Submissions (  
  record_id integer UNIQUE,  
  assignment_id integer,  
  student_id integer,  
  time_submitted integer,  
  grade_received byte,  
  
  PRIMARY KEY(assignment_id, student_id)  
);
```

Assume the column record\_id corresponds to the row's actual record ID.

Q1: How large (in bytes) is a record?

We simply add up the sizes of each field in a record. We have 4 integer fields and 1 byte field, which is  $4 \times 4 = 17$  bytes.

Q2: Suppose we begin each page with a 32-bytes header plus a bitmap. At most, how many records can fit in an unpacked page?

Let's convert everything into bits. First, a page holds  $1,024 \times 64 \times 8 = 1,048,576$  bits while a record holds  $17 \times 8 = 136$  bits. Now, remember that in an unpacked page, each record needs an additional bit to represent whether or not it is valid (e.g. has been deleted). This means we need 1 more bit per record, so each record in fact requires 137 bits. Finally, we have an extra 32 bytes ( $32 \times 8 = 256$  bits) reserved for the page header. This gives us a maximum of  $(1,048,576 - 256) / 137 \approx 7,651$  records.

We add two variable-length fields to our table schema. Now our table looks like this:

```
CREATE TABLE Submissions (  
  record_id integer UNIQUE,  
  assignment_id integer,  
  student_id integer,  
  time_submitted integer,  
  grade_received byte,  
  
  comment text,  
  regrade_request text,  
  
  PRIMARY KEY(assignment_id, student_id)  
);
```

We decide to use slotted pages to store the variable length records. Each page begins with a 24-byte header plus a slot directory. (Assume this header contains information such as the number of valid records in the page.) Each pointer inside the slot directory consumes 20 bits/record, while the record header storing field offsets is 32 bits wide.

Q3: What is the maximum number of records that can fit in our slotted pages?

Again, we have 1,048,576 bits per page, the page header consumes 256 bits, and a record is 136 bits wide. Now instead of a bitmap, where each record takes 1 extra bits, we have a slot directory of pointers, so each record requires 20 extra bits. Thus each record costs us  $136 + 20 = 156$  bits. We need to store field offsets for the variable-length text fields, which is an additional 32 bits per record, bringing us to a total of  $156 + 32 = 188$  bits per record. Finally, note that we get the maximum possible number of records when all comment and regrade\_request fields are NULL; i.e. they both take up 0 bytes. Then the smallest possible memory footprint per record is 188 bits/record, so we get  $(1,048,576 - 256) / 188 \approx 5,576$  records.

Q4: We decide to squash the two text fields together into one field using a semicolon separator character (;), which allows us to shrink the record header from 32 bits to 16 bits at the cost of 8 bits (for the semicolon). For example, the columns ("Submitted late", "Dog ate my homework") get compressed into "Submitted late;Dog ate my homework". Which of the following are true with this new scheme?

Since we're using the semicolon as a field separator, we can't enter comments with semicolons. In our example, the comment "Fantastic work; good job!" would be truncated to "Fantastic work" and " good job!" would be misinterpreted as the regrade\_request field. However,

since our records become 8 bits shorter, we'll be able to fit more records per page. As a result, our table will become smaller, and table scans will speed up accordingly (depending on how many pages we save on storage).

**C, D**