

CS150A Database

Lu Sun

School of Information Science and Technology

ShanghaiTech University

Sept. 8, 2022

Today:

- DML in Multiple Tables
 - Set Operations
 - Nested Queries
 - Join
- Null Values

Readings:

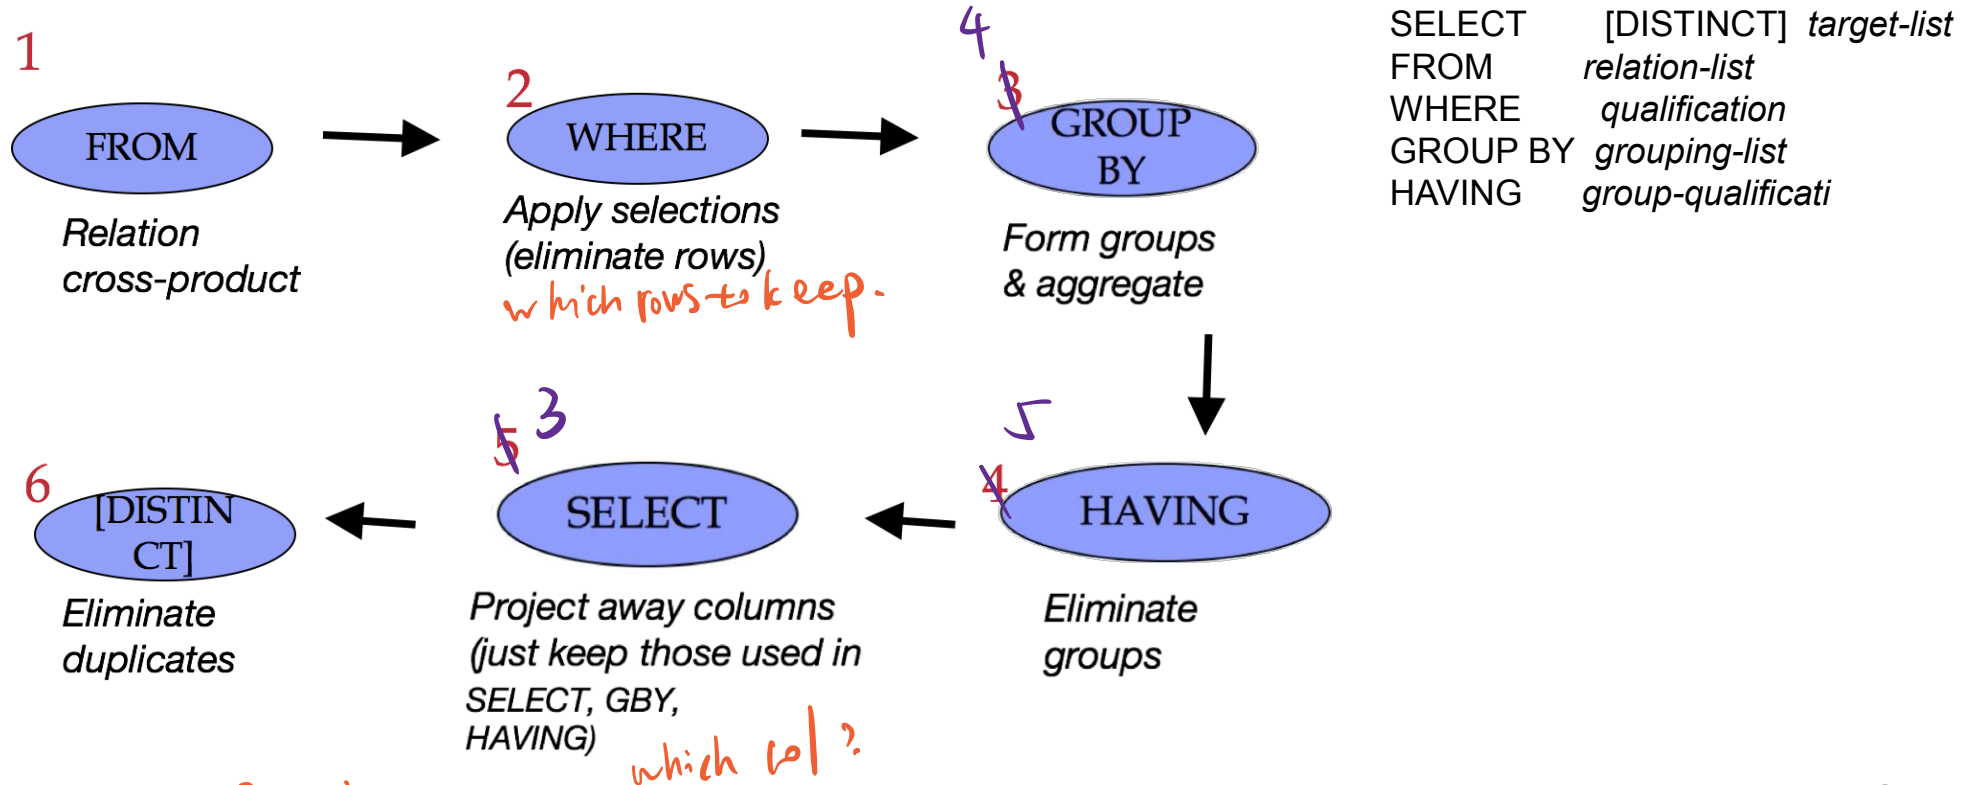
- Database Management Systems (DBMS), Chapter 5
- Lecture note SQL II

SQL DML 1:

Basic Single-Table Queries

- **SELECT** [**DISTINCT**] *<column expression list>*
FROM *<single table>*
[**WHERE** *<predicate>*]
[**GROUP BY** *<column list>*
[**HAVING** *<predicate>*]]
[**ORDER BY** *<column list>*]
[**LIMIT** *<integer>*];

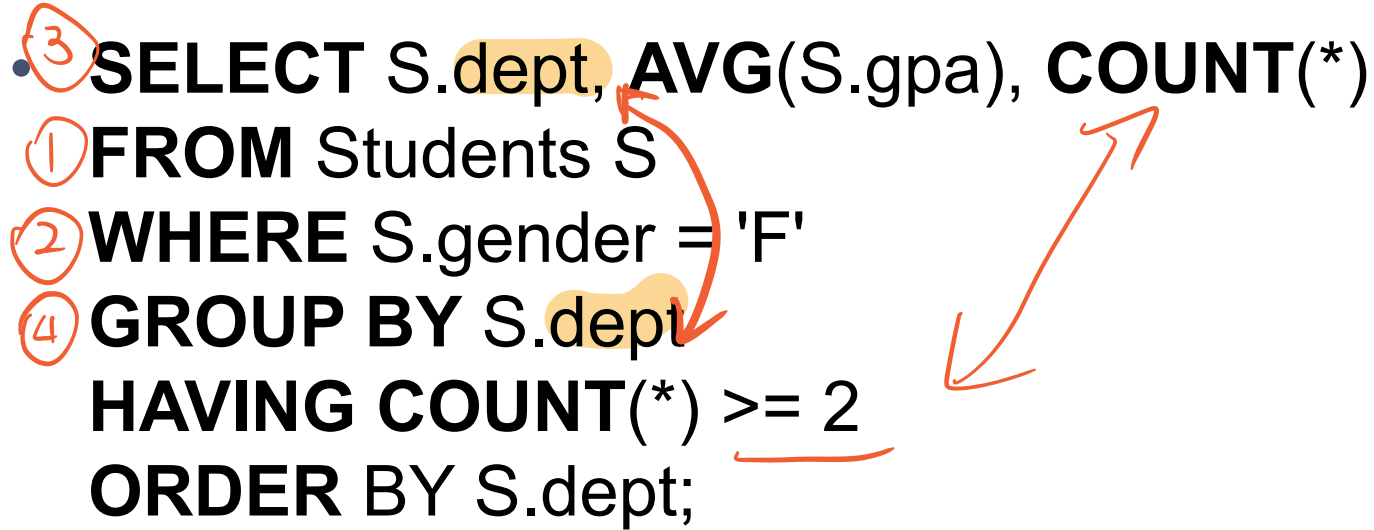
Conceptual SQL Evaluation



And Order By, Limit ...

Putting it all together

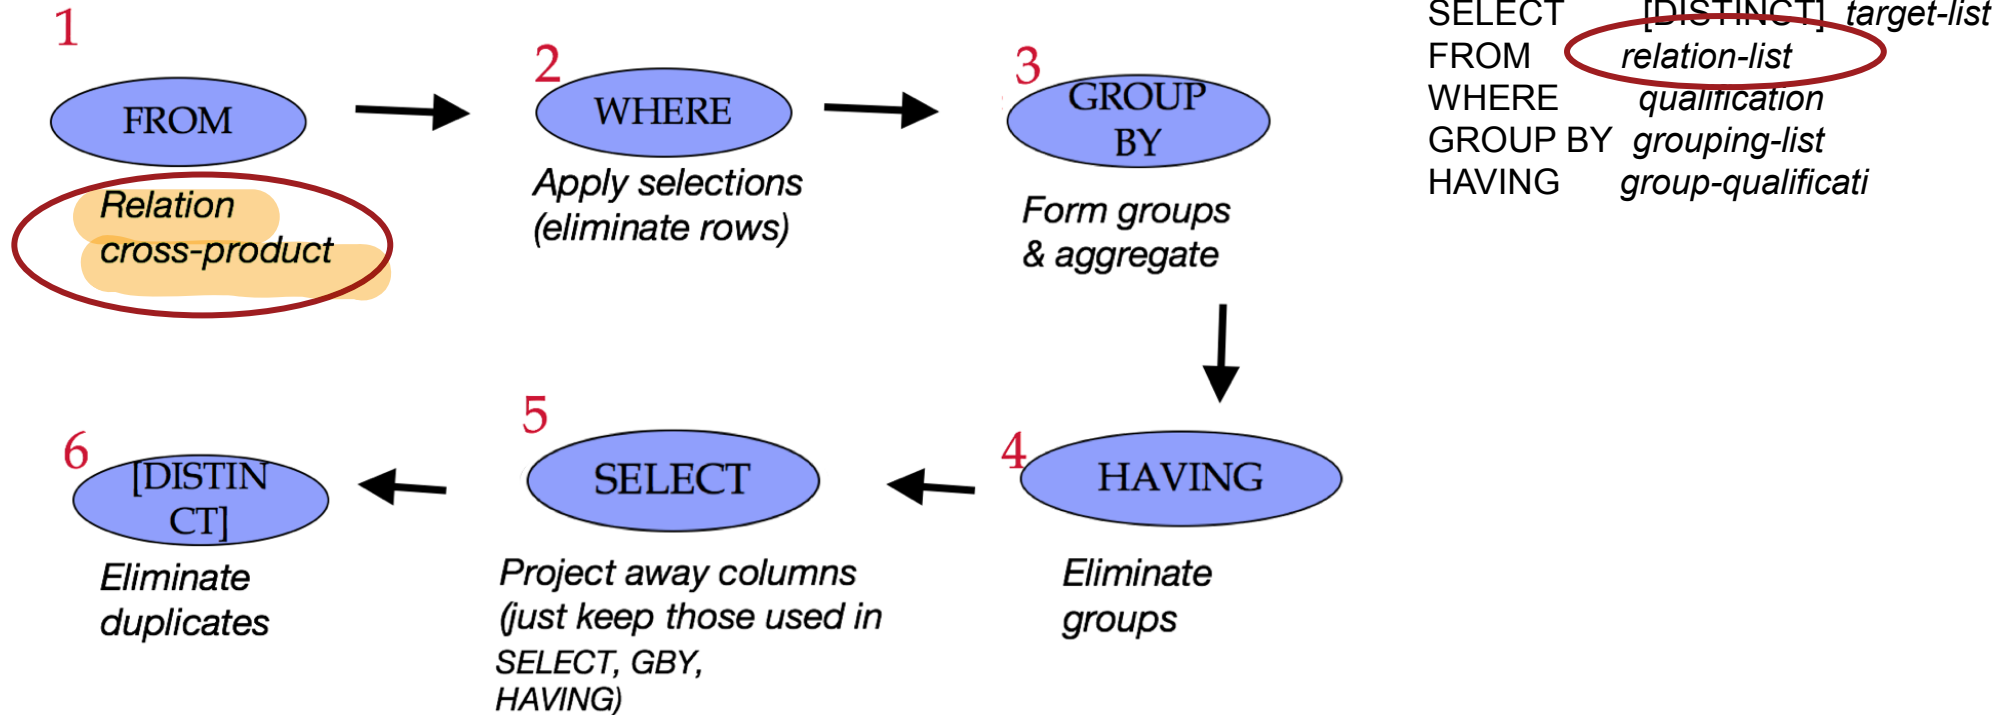
• ③ **SELECT** S.dept, **AVG**(S.gpa), **COUNT**(*)
① **FROM** Students S
② **WHERE** S.gender = 'F'
④ **GROUP BY** S.dept
HAVING COUNT(*) >= 2
ORDER BY S.dept;



Join Queries

- SELECT [DISTINCT] *<column expression list>*
FROM *<table1 [AS t1], ... , tableN [AS tn]>*
[WHERE *<predicate>*]
[GROUP BY *<column list>*[HAVING *<predicate>*]]
[ORDER BY *<column list>*];

Conceptual SQL Evaluation, cont



Cross (Cartesian) Product


- All pairs of tuples, concatenated

Sailors

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

Reserves

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01



sid	sname	rating	age	sid	bid	day
1	Popeye	10	22	1	102	9/12
1	Popeye	10	22	2	102	9/13
1	Popeye	10	22	1	101	10/01
2	OliveOyl	11	39	1	102	9/12
...

Find sailors who've reserved a boat

```
SELECT S.sid, S.sname, R.bid
FROM Sailors AS S, Reserves AS R
WHERE S.sid=R.sid
```

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

sid	sname	rating	age	sid	bid	day
1	Popeye	10	22	1	102	9/12
1	Popeye	10	22	2	102	9/13
1	Popeye	10	22	1	101	10/01
2	OliveOyl	11	39	1	102	9/12
...

Find sailors who've reserved a boat cont

```
SELECT S.sid, S.sname, R.bid  
FROM Sailors AS S, Reserves AS R  
WHERE S.sid=R.sid
```

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

sid	sname	bid
1	Popeye	102
1	Popeye	101
2	OliveOyl	102

Column Names and Table Aliases

Range var

Sailors.sname Ok

```
SELECT Sailors.sid, sname, bid  
FROM Sailors, Reserves  
WHERE Sailors.sid = Reserves.sid
```

```
SELECT S.sid, sname, bid  
FROM Sailors AS S, Reserves AS R  
WHERE S.sid = R.sid
```

Alias.

Self join

More Aliases

```
SELECT x.sname, x.age,  
       y.sname AS sname2,  
       y.age AS age2  
FROM Sailors AS x, Sailors AS y  
WHERE x.age > y.age
```

renaming output

sname	age	sname2	age2
Popeye	22	Bob	19
OliveOyl	39	Popeye	22
OliveOyl	39	Garfield	27
OliveOyl	39	Bob	19
Garfield	27	Popeye	22
Garfield	27	Bob	19

- Table aliases in the FROM clause
 - Needed when the same table used multiple times (“self-join”)
- Column aliases in the SELECT clause

Arithmetic Expressions

- ```
SELECT S.age, S.age-5 AS age1, 2*S.age AS age2
FROM Sailors AS S
WHERE S.sname = 'Popeye'
```
- ```
SELECT S1.sname AS name1, S2.sname AS name2
FROM   Sailors AS S1, Sailors AS S2
WHERE  2*S1.rating = S2.rating - 1
```

SQL Calculator!

SELECT

log(1000) as three,

exp(ln(2)) as two,

cos(0) as one,

ln(2*3) = ln(2) + ln(3) as sanity;

three two one sanity
 $\log(1000)$ $\exp(\ln 2)$ $\cos(0)$ $\ln(2 \cdot 3) = \ln 2 + \ln 3$

String Comparisons

- Old School SQL

```
SELECT S.sname  
FROM   Sailors S  
WHERE  S.sname LIKE 'B_%'
```

B. ... 姓名以B开头的

- Standard **Regular Expressions**

```
SELECT S.sname  
FROM   Sailors S  
WHERE  S.sname ~ 'B.*'
```

B.*

Combining Predicates

- Subtle connections between:
 - Boolean logic in WHERE (i.e., AND, OR)
 - Traditional Set operations (i.e. INTERSECT, UNION)
- Let's see some examples...

Sid's of sailors who reserved a red **OR** a green boat

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND  
      (B.color='red' OR B.color='green')
```


Sid's of sailors who reserved a red **OR** a green boat Pt 2

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND  
      (B.color='red' OR B.color='green')
```

VS...

|| eq

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'
```

UNION ALL

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='green'
```

Sid's of sailors who reserved a red **AND** a green boat Pt 3

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND  
      (B.color='red' AND B.color='green')
```

→ Nothing


~~Not possible~~

VS...

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'
```

INTERSECT

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='green'
```



As the diagram shows, when you intersect query 1 (green) and query 2 (blue), the result will be the intersection of the two, the grey color in the above diagram.

Find sailors who have **not** reserved a boat

Except

```
SELECT S.sid  
FROM   Sailors S
```

EXCEPT

```
SELECT S.sid  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid
```

Set Semantics

- Set: a collection of *distinct* elements
- Standard ways of manipulating/combining sets
 - Union
 - Intersect
 - Except
- Treat tuples within a relation as elements of a set

Default: Set Semantics

Note: R and S are **relations**. They are not sets, since they have duplicates.

$R = \{A, A, A, A, B, B, C, D\}$

$S = \{A, A, B, B, B, C, E\}$

- UNION

$\{A, B, C, D, E\}$

- INTERSECT

$\{A, B, C\}$

- EXCEPT

$\{D\}$

*remove
duplicate*

Note: Think of each letter as being a **tuple** in **relation**.

ex:

A: (Jim, 18, English, 4.0)

B: (Marcela, 20, CS, 3.8)

C: (Gail, 19, Statistics, 3.74)

D: (Goddard, 20, Math, 3.8)

“ALL” : Multiset Semantics

$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

“UNION ALL” : Multiset Semantics

$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

- UNION ALL: sum of cardinalities

$\{A(4+2), B(2+3), C(1+1), D(1+0), E(0+1)\}$

$= \{A, A, A, A, A, A, B, B, B, B, B, C, C, D, E\}$

“INTERSECT ALL” : Multiset Semantics

$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

- INTERSECT ALL: min of cardinalities
 $\{A(\min(4,2)), B(\min(2,3)), C(\min(1,1)),$
 $D(\min(1,0)), E(\min(0,1))\}$
 $= \{A, A, B, B, C\}$

“EXCEPT ALL” : Multiset Semantics

$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

- EXCEPT ALL: difference of cardinalities
 $\{A(4-2), B(2-3), C(1-1), D(1-0), E(0-1)\}$
 $= \{A, A, D, \}$


Nested Queries: IN



- *Names of sailors who've reserved boat #102:*

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid IN
      (SELECT R.sid
       FROM   Reserves R
       WHERE  R.bid=102)
```

subquery



Nested Queries: NOT IN

- *Names of sailors who've not reserved boat #103:*

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid NOT IN
        (SELECT  R.sid
         FROM    Reserves R
         WHERE   R.bid=103)
```

Nested Queries: EXISTS

- *This is a bit odd, but it is legal:*

```
SELECT  S.sname  
FROM    Sailors S  
WHERE   EXISTS  
        (SELECT  R.sid  
         FROM    Reserves R  
         WHERE   R.bid=103)
```

if exist R.bid
Sailor

return all
sailor

or, return NaN

Nested Queries with Correlation

- Names of sailors who've reserved boat #102:

```
SELECT  S.sname
FROM    Sailors S
WHERE   EXISTS
        (SELECT *
         FROM Reserves R
         WHERE R.bid=102 AND S.sid=R.sid)
```

correlation

- Correlated subquery is recomputed for each Sailors tuple.

execute for every sailor tuple.

More on Set-Comparison Operators

- We've seen: IN, EXISTS
- Can also have: NOT IN, NOT EXISTS
- Other forms: op ANY, op ALL

Find sailors whose rating is greater than that of **some sailor** called Popeye:

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating > ANY  
      (SELECT S2.rating  
       FROM   Sailors S2  
       WHERE  S2.sname='Popeye')
```

*maybe
multiple.*

A Tough One: “Division”

- Relational Division: “Find sailors who’ve reserved all boats.”
Said differently: “sailors with no counterexample missing boats”

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
  (SELECT B.bid
   FROM Boats B
   WHERE NOT EXISTS (SELECT R.bid
                     FROM Reserves R
                     WHERE R.bid=B.bid
                           AND R.sid=S.sid ))
```

nested 2: for R
(3rd Table)

ARGMAX? Pt 1

- The sailor with the highest rating
- Correct or Incorrect?

Just max, only outputs the rating

```
SELECT MAX(S.rating)
FROM Sailors S;
```

VS

still

```
SELECT S.*, MAX(S.rating)
FROM Sailors S;
```

No Groupy here, so by default group all

illegal

*"groupby - select"
(col)
Not match*

only when each group has a

ARGMAX? Pt 2

legal value for the selected col

- The sailor with the highest rating
- Correct or Incorrect? Same or different?

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating >= ALL  
      (SELECT S2.rating  
       FROM   Sailors S2)
```

✓

legal, same

VS

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating =  
      (SELECT MAX(S2.rating)  
       FROM   Sailors S2)
```

✓

agg max in nested

ARGMAX? Pt 3

- The sailor with the highest rating
- Correct or Incorrect? Same or different?

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating >= ALL  
      (SELECT S2.rating  
       FROM   Sailors S2)
```

→ all the sailors.

VS

```
SELECT *  
FROM   Sailors S  
ORDER BY rating DESC  
LIMIT 1;
```

not deterministic

↑ pick arbitrary from
max rating sailors.

“Inner” Joins: Another Syntax

Normal Join:

```
SELECT s.*, r.bid  
FROM Sailors s, Reserves r  
WHERE s.sid = r.sid  
AND ...
```

Inner Join:

```
SELECT s.*, r.bid  
FROM Sailors s INNER JOIN Reserves r  
ON s.sid = r.sid  
WHERE ...
```

|| eq.

Join Variants

```
SELECT <column expression list>
FROM table_name
  [INNER | NATURAL
   | {LEFT | RIGHT | FULL } {OUTER}] JOIN
  table_name
ON <qualification_list>
WHERE ...
```

- INNER is default
- Inner join what we've learned so far
 - Same thing, just with different syntax.

Inner/Natural Joins

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors s, Reserves r  
WHERE s.sid = r.sid  
AND s.age > 20;
```

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors s INNER JOIN Reserves r  
ON s.sid = r.sid  
WHERE s.age > 20;
```

(Not good)

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors s NATURAL JOIN Reserves r  
WHERE s.age > 20;
```

eq

on auto (matching col name)

- **ALL 3 ARE EQUIVALENT!**
- “NATURAL” means equi-join for pairs of attributes with the same name

Left Outer Join

- Returns all matched rows, and preserves all unmatched rows from the table on the left of the join clause
 - (use nulls in fields of non-matching tuples)

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors2 s LEFT OUTER JOIN Reserves2 r  
ON s.sid = r.sid;
```

Returns all sailors & bid for boat in any
of their reservations

Note: no match for s.sid? r.bid IS NULL!

SQL Script

```

1 CREATE TABLE TS (
2   id INT(6) UNSIGNED AUTO_INCREMENT PRIMARY KEY,
3   firstname VARCHAR(30) NOT NULL,
4   lastname VARCHAR(30) NOT NULL,
5   email VARCHAR(50),
6   reg_date TIMESTAMP
7 );
8
9 CREATE TABLE TB (
10  bid INT(6) UNSIGNED AUTO_INCREMENT PRIMARY KEY,
11  firstname VARCHAR(30) NOT NULL,
12  lastname VARCHAR(30) NOT NULL,
13  reg_date TIMESTAMP
14 );
15
16
17 INSERT INTO 'TS' ('id', 'firstname', 'lastname', 'email', 'reg_date') VALUES ('1', 'John', 'Doe', 'john.doe@sqltest.net', CURRENT_TIMESTAMP);
18 INSERT INTO 'TS' ('id', 'firstname', 'lastname', 'email', 'reg_date') VALUES ('21', 'J2ohn', 'D2oe', 'j2ohn.doe@sqltest.net', CURRENT_TIMESTAMP);
19 INSERT INTO 'TB' ('bid', 'firstname', 'lastname', 'reg_date') VALUES ('21', 'J2ohn', 'D2oe', CURRENT_TIMESTAMP);

```

Load Test Example

Clear

SQL Query

```

1 SELECT TS.id, TS.firstname, TB.bid
2 FROM TS LEFT OUTER JOIN TB
3 ON TS.id = TB.bid

```

Click on a [Execute SQL](#) button

Všeč mi je



Děli z drugimi

Link for sharing your example: <https://sqltest.net/#1879029>

Result

id	firstname	bid
1	John	Null, (Not match)
21	J2ohn	21

SQL Script

```

1 CREATE TABLE TS (
2   id INT(6) UNSIGNED AUTO_INCREMENT PRIMARY KEY,
3   firstname VARCHAR(30) NOT NULL,
4   lastname VARCHAR(30) NOT NULL,
5   email VARCHAR(50),
6   reg_date TIMESTAMP
7 );
8
9 CREATE TABLE TB (
10  bid INT(6) UNSIGNED AUTO_INCREMENT PRIMARY KEY,
11  firstname VARCHAR(30) NOT NULL,
12  lastname VARCHAR(30) NOT NULL,
13  reg_date TIMESTAMP
14 );
15
16
17 INSERT INTO 'TS' ('id', 'firstname', 'lastname', 'email', 'reg_date') VALUES ('1', 'John', 'Doe', 'john.doe@sqltest.net', CURRENT_TIMESTAMP);
18 INSERT INTO 'TS' ('id', 'firstname', 'lastname', 'email', 'reg_date') VALUES ('21', 'J2ohn', 'D2oe', 'j2ohn.doe@sqltest.net', CURRENT_TIMESTAMP);
19 INSERT INTO 'TB' ('bid', 'firstname', 'lastname', 'reg_date') VALUES ('21', 'J2ohn', 'D2oe', CURRENT_TIMESTAMP);

```

Load Test Example

Clear

SQL Query

```

1 SELECT TS.id, TS.firstname, TB.bid
2 FROM TS INNER JOIN TB
3 ON TS.id = TB.bid

```

Click on a [Execute SQL](#) button

Všeč mi je



Děli z drugimi

Link for sharing your example: <https://sqltest.net/#1879030>

Result

id	firstname	bid
21	J2ohn	21

preserve tuple (selected cols)
 doesn't match in Left
 side Table.

Right Outer Join

- Returns all matched rows, and preserves all unmatched rows from the table on the right of the join clause
 - (use nulls in fields of non-matching tuples)

```
SELECT r.sid, b.bid, b.bname  
FROM Reserves2 r RIGHT OUTER JOIN Boats2 b  
ON r.bid = b.bid
```

Returns all boats and sid for any sailor associated with the reservation.

Note: no match for b.bid? r.sid IS NULL!

Full Outer Join

both side Not match
preserve.

- Returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.bname  
FROM Reserves2 r FULL OUTER JOIN Boats2 b  
ON r.bid = b.bid
```

- Returns all boats & all information on reservations
- No match for r.bid?
 - b.bid IS NULL AND b.bname IS NULL!
- No match for b.bid?
 - r.sid IS NULL!

Views: Named Queries

```
CREATE VIEW view_name  
AS select_statement
```

- Makes development simpler
- Often used for security
- Not “materialized”

```
CREATE VIEW Redcount  
AS SELECT B.bid, COUNT(*) AS scount  
   FROM Boats2 B, Reserves2 R  
  WHERE R.bid=B.bid AND B.color='red'  
  GROUP BY B.bid
```

cont Num of reservations for each red boat.


Quirey's Name.

Views Instead of Relations in Queries

```
CREATE VIEW Redcount
AS SELECT B.bid, COUNT(*) AS scount
   FROM Boats2 B, Reserves2 R
   WHERE R.bid=B.bid AND B.color='red'
   GROUP BY B.bid;
```

all

```
SELECT * from redcount;
```



```
SELECT bname, scout
FROM Redcount R, Boats2 B
WHERE R.bid=B.bid
AND scout < 10;
```

bid	scount
102	1

Subqueries in FROM

→ delete auto later.

Like a “view on the fly”!

```
SELECT bname, scout  
FROM Boats2 B,
```

‘Red count’


```
(SELECT B.bid, COUNT (*)  
  FROM Boats2 B, Reserves2 R  
 WHERE R.bid = B.bid AND B.color = 'red'  
 GROUP BY B.bid) AS Reds(bid, scout)
```

```
WHERE Reds.bid=B.bid  
      AND scout < 10
```

```
SELECT bname, scount  
FROM Redcount R, Boats2 B  
WHERE R.bid=B.bid  
AND scount < 10;
```

WITH a. k. a. common table expression (CTE)

Another “view on the fly” syntax:



```
WITH Reds(bid, scout) AS  
(SELECT B.bid, COUNT (*)  
FROM Boats2 B, Reserves2 R  
WHERE R.bid = B.bid AND B.color = 'red'  
GROUP BY B.bid)
```

```
SELECT bname, scout  
FROM Boats2 B, Reds  
WHERE Reds.bid=B.bid  
AND scout < 10
```

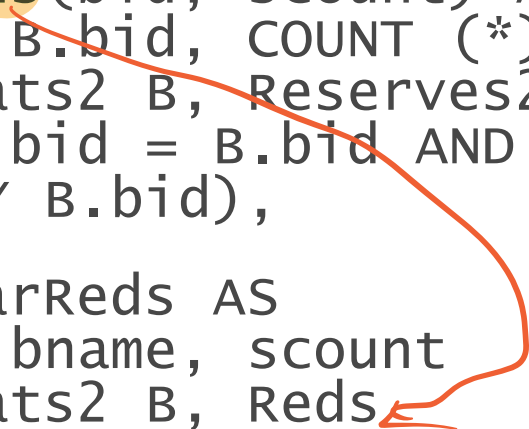
```
SELECT bname, scout  
FROM Redcount R, Boats2 B  
WHERE R.bid=B.bid  
AND scout < 10;
```

Can have many queries in WITH

Another “view on the fly” syntax:

```
WITH Reds (bid, scout) AS  
  (SELECT B.bid, COUNT (*)  
   FROM Boats2 B, Reserves2 R  
   WHERE R.bid = B.bid AND B.color = 'red'  
   GROUP BY B.bid),
```

```
UnpopularReds AS  
  (SELECT bname, scout  
   FROM Boats2 B, Reds  
   WHERE Reds.bid=B.bid  
   AND scout < 10)
```



```
SELECT * FROM UnpopularReds;
```

ARGMAX GROUP BY?

- 1 • The sailor with the highest rating per age 2 3

WITH maxratings(age, maxrating) AS
(SELECT age, max(rating)
FROM Sailors
GROUP BY age)

(2, 3)

SELECT S.*
FROM Sailors S, maxratings m
WHERE S.age = m.age
AND S.rating = m.maxrating;

1, (2, 3)



Brief Detour: Null Values

- Field values are sometimes unknown
 - SQL provides a special value NULL for such situations.
 - Every data type can be NULL
- The presence of null complicates many issues. E.g.:
 - Selection predicates (WHERE)
 - Aggregation
- But NULLs comes naturally from Outer joins

NULL in the WHERE clause

- Consider a tuple where `rating` IS NULL.

```
INSERT INTO sailors VALUES  
(11, 'Jack Sparrow', NULL, 35);
```

```
SELECT * FROM sailors  
WHERE rating > 8;
```

Is Jack Sparrow in the output?

NULL in comparators

- Rule: (x op NULL) evaluates to ... NULL!

```
SELECT 100 = NULL;
```

```
SELECT 100 < NULL;
```

```
SELECT 100 >= NULL;
```

Explicit NULL Checks

opt for Null

```
SELECT * FROM sailors WHERE rating IS NULL;
```

```
SELECT * FROM sailors WHERE rating IS NOT NULL;
```

NULL at top of WHERE

- Rule: Do not output a tuple WHERE NULL

SELECT * FROM sailors;

(null > 8) → Null

SELECT * FROM sailors WHERE rating > 8;

SELECT * FROM sailors WHERE rating <= 8;

Not output

NULL in Boolean Logic

Three-valued logic:

NOT	T	F	N
	F	T	

AND	T	F	N
T	T	F	
F	F	F	
N			

OR	T	F	N
T	T	T	
F	T	F	
N			

`SELECT * FROM sailors WHERE rating > 8 AND TRUE;`

`SELECT * FROM sailors WHERE rating > 8 OR TRUE;`

`SELECT * FROM sailors WHERE NOT (rating > 8);`

General rule: NULL can take on either 'TRUE' or 'FALSE', so answers need to accommodate either value.

NULL in Boolean Logic

Three-valued logic:

NOT	T	F	N
	F	T	N

AND	T	F	N
T	T	F	N
F	F	F	F
N	N	F	N

OR	T	F	N
T	T	T	T
F	T	F	N
N	T	N	N

`SELECT * FROM sailors WHERE rating > 8 AND TRUE;`

`SELECT * FROM sailors WHERE rating > 8 OR TRUE;`

`SELECT * FROM sailors WHERE NOT (rating > 8);`

General rule: NULL can take on either 'TRUE' or 'FALSE', so answers need to accommodate either value.

NULL and Aggregation

only cover
SELECT **count(*)** FROM sailors;

Cover Null

SELECT count(rating) FROM sailors;

Not cover Null

SELECT sum(rating) FROM sailors;

Not cover

SELECT avg(rating) FROM sailors;

Not cover

General rule: NULL **column values
are ignored by aggregate functions**

NULLs: Summary

- **x op NULL is NULL**
- WHERE NULL: do not send to output
- Boolean connectives: 3-valued logic
- Aggregates ignore NULL-valued inputs

Testing SQL Queries

- SQL Fiddle pages <http://sqlfiddle.com/> will typically help you answer the questions in the worksheets and quizzes.
- But in real life:
 - not every database instance will reveal every bug in your query.
 - Eg: database instance without any rows in it!
 - Need to debug your queries
 - reasoning about them carefully
 - constructing test data.

Tips for Generating Test Data

- Generate **random data**
 - e.g. using a service like <https://mockaroo.com/>
- Try to construct data that could check for the following potential errors:
 - Incorrect output **schema**
 - Output may be **missing rows** from the correct answer (false negatives)
 - Output may contain **incorrect rows** (false positives)
 - Output may have the **wrong number of duplicates**.
 - Output may **not be ordered properly**.

Summary

- You've now seen SQL—you are armed.
- A declarative language
 - Somebody has to translate to algorithms though...
 - The RDBMS implementer ... i.e. you!

Summary Cont

- The data structures and algorithms that make SQL possible also power:
 - NoSQL, data mining, scalable ML, network routing...
 - A toolbox for scalable computing!
 - That fun begins next week
- We skirted questions of good database (schema) design
 - a topic we'll consider in greater depth later