

CS150A Database

Lu Sun

School of Information Science and Technology

ShanghaiTech University

Sept. 20, 2022

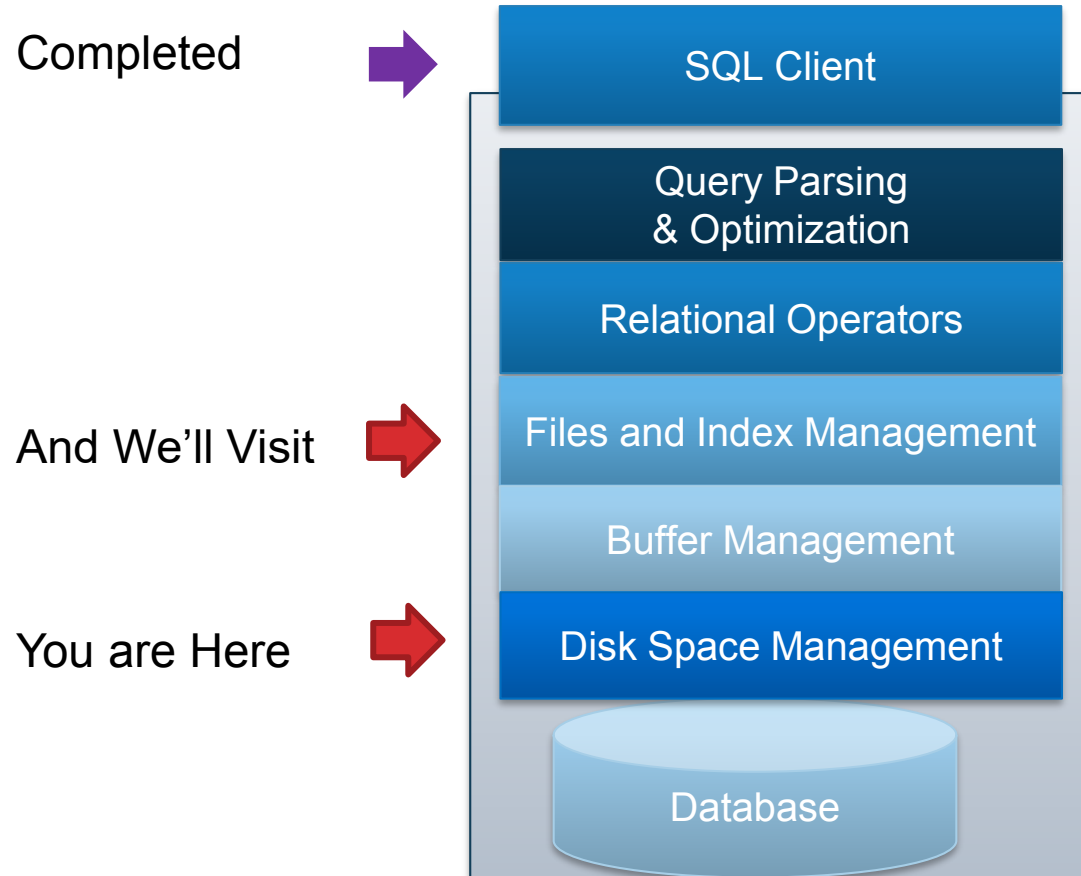
Today:

- Cost Model and Analysis:
 - In Heap File and Sorted File

Readings:

- Database Management Systems (DBMS), Chapters 8 and 9

Architecture of a DBMS



Recall: Heap Files

- Unordered **collection** of records
- Recall API for higher layers of the DBMS.
Today we'll ask: "How? At what cost?"
 - Insert/delete/modify record
 - Fetch a particular record by ***record id*** ...
 - Record id is a pointer encoding pair of (**pageID**, **location** on page)
 - Scan all records
 - Possibly with some conditions on the records to be retrieved

Recall: Multiple File Organizations

Many alternatives exist, each good in some situations and less so in others.

This is a theme in DB systems work!

- **Heap Files:** Suitable when typical access is a full scan of all records
- **Sorted Files:** Best for retrieval in order, or when a range of records is needed
- **Clustered Files & Indexes:** Group data into blocks to enable fast lookup *and* efficient modifications.
 - More on this soon ...

Bigger Questions

- What is the “best” file organization?
 - Depends on access patterns ...
 - How? What are common access patterns anyway?
- Can we be quantitative about tradeoffs?
 - If one is better ... by how much?

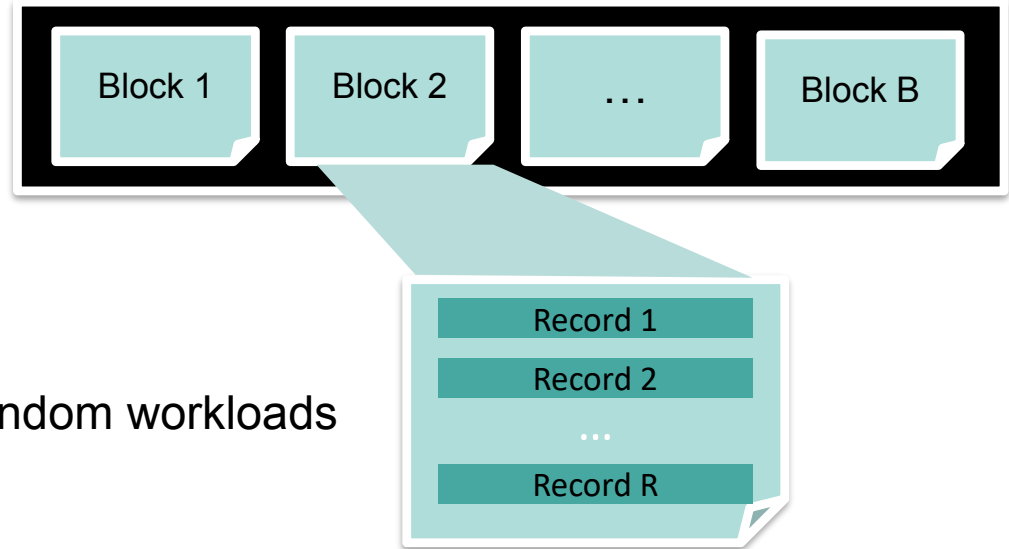
Goals

- Big picture overheads for data access
 - We'll (overly) simplify performance models to provide insight, not to get perfect performance
 - Still, a bit of discipline:
 - **Clearly identify assumptions up front**
 - **Then estimate cost in a principled way**
- Foundation for query optimization
 - Can't choose the fastest scheme without an estimate of speed!

COST MODEL AND ANALYSIS

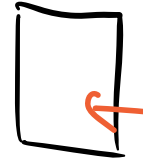
Cost Model for Analysis

- **B**: The number of data blocks in the file
- **R**: Number of records per block
- **D**: (Average) time to read/write disk block
- Focus: Average case analysis for uniform random workloads
- For now, we will ignore
 - Sequential vs Random I/O
 - Pre-fetching
 - Any in-memory costs
- Good enough to show the overall trends



More Assumptions

- **Single record** insert and delete
- Equality selection – **exactly one match**
- For **Heap Files**:
 - Insert always **appends to end of file.**
- For Sorted Files:
 - **Packed**: Files compacted after deletions.
 - Sorted according to search key



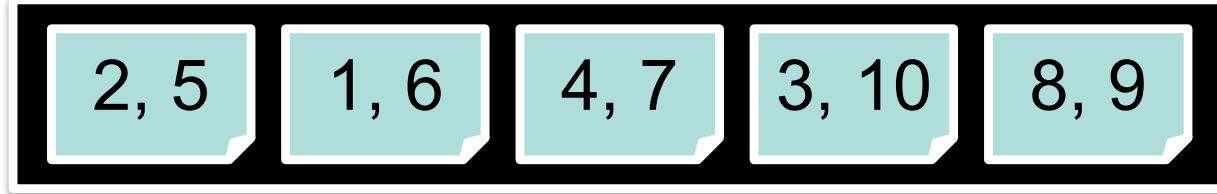
no holes.

Extra Challenge

- After understanding these slides ...
 - You should question all these assumptions and rework
 - Good exercise to study for tests, and generate ideas

Heap Files & Sorted Files

Heap File



Sorted File



For illustration, records are just integers

- **B:** The number of data blocks = 5
- **R:** Number of records per block = 2
- **D:** (Average) time to read/write disk block = 5ms

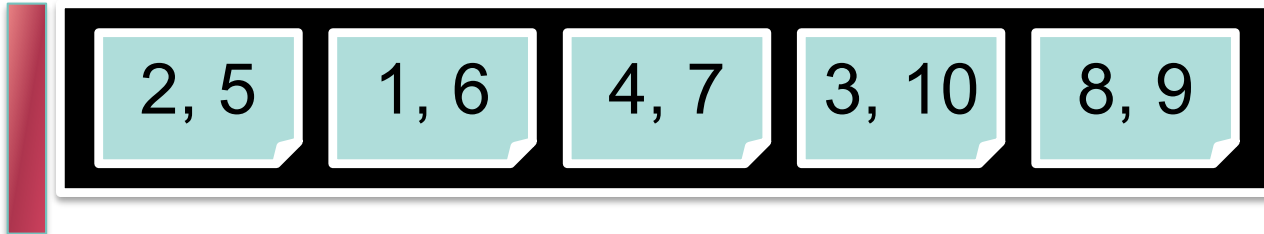
Cost of Operations: Scan?

	Heap File	Sorted File
Scan all records	↙	
Equality Search		
Range Search		
Insert		
Delete		

- **B:** The number of data blocks = 5
- **R:** Number of records per block = 2
- **D:** (Average) time to read/write disk block = 5ms

Scan All Records

Heap File



Sorted File



- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk block
- **Pages touched:** ?
- **Time to read the record:** ?

Cost of Operations: Scan Cost

	Heap File	Sorted File
Scan all records	$B \cdot D$	$B \cdot D$
Equality Search		
Range Search		
Insert		
Delete		

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk bloc

Cost of Operations: Equality Search?

	Heap File	Sorted File
Scan all records	$B \cdot D$	$B \cdot D$
Equality Search		
Range Search		
Insert		
Delete		

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk block

Find Key 8: Heap File

Heap File



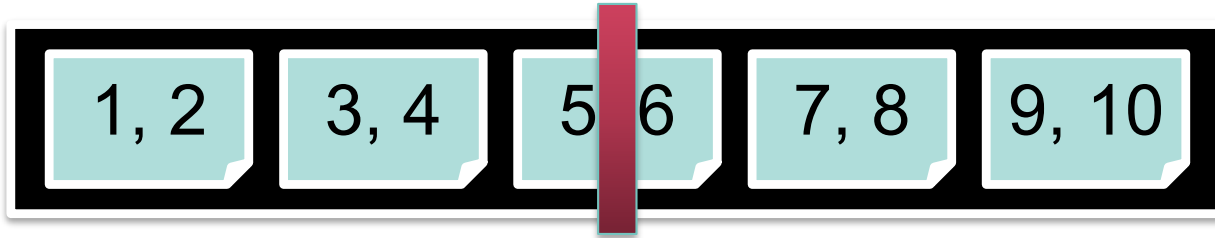
- **P(i)**: Probability that key is on page i is $1/B$
- **T(i)**: Number of pages touched if key on page i is i
- Therefore the expected number of pages touched
- **Pages touched on average?**

calc

$$\sum_{i=1}^B T(i)P(i) = \sum_{i=1}^B i \frac{1}{B} = \frac{B(B+1)}{2B} \approx \frac{B}{2}$$

Find Key 8: Sorted File

Sorted File



binary search

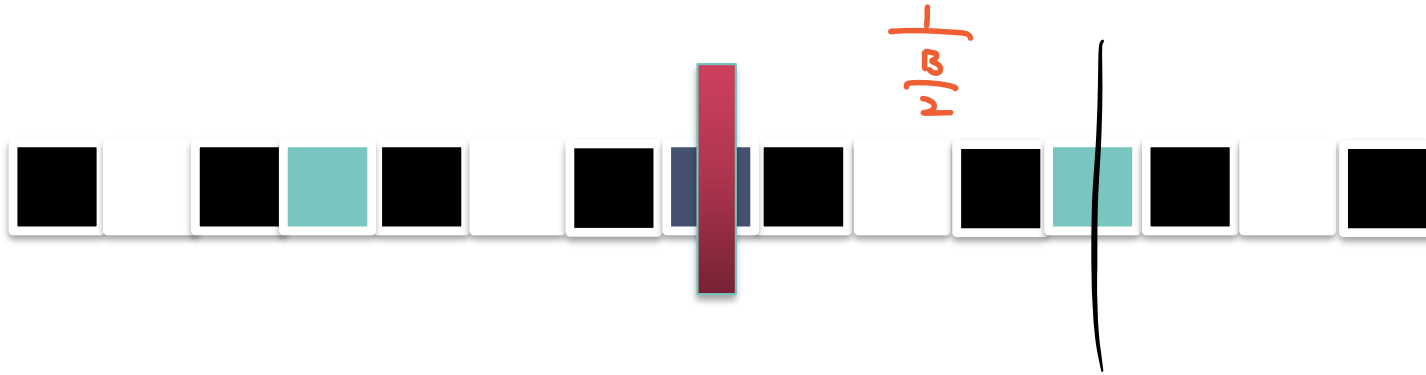
- **Worst-case:** Pages touched in binary search
 - $\log_2 B$
- **Average-case:** Pages touched in binary search
 - $\log_2 B?$

Average Case Binary Search

read time

read 3 time
to get
to this level

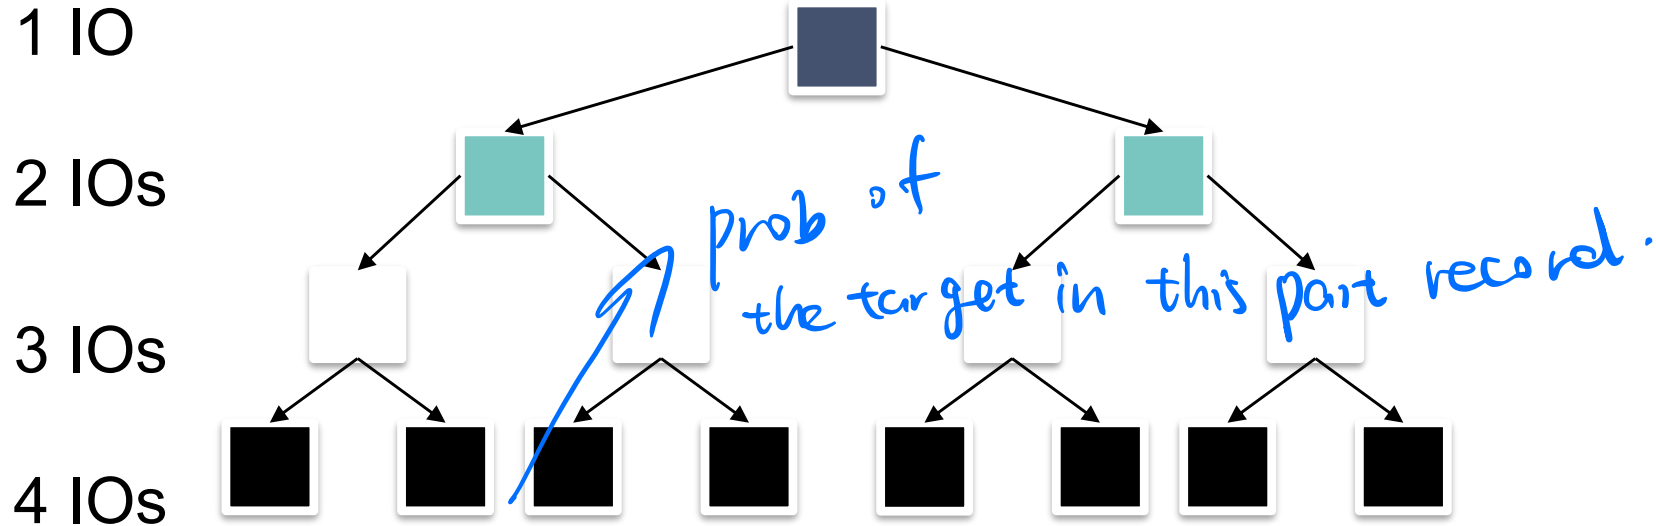
Expected Number of Reads: $1(1/B) + 2(2/B) + 3(4/B) + 4(8/B)$



Average Case Binary Search cont

↓ prob of answer in middle page

Expected Number of Reads: $1 (1 / B) + 2 (2 / B) + 3 (4 / B) + 4 (8 / B)$



$$\sum_{i=1}^{\log_2 B} \frac{i 2^{i-1}}{B} = \frac{1}{B} \sum_{i=1}^{\log_2 B} i 2^{i-1} = \log_2 B - \frac{B-1}{B} = \log_2 B.$$

lost to get the level (read times)

Cost of Operations: Equation Search Cost

	Heap File	Sorted File
Scan all records	$B \cdot D$	$B \cdot D$
Equality Search	$0.5 \cdot B \cdot D$	$(\log_2 B) \cdot D$
Range Search		
Insert		
Delete		

- **B**: The number of data blocks
- **R**: Number of records per block
- **D**: Average time to read/write disk block

Cost of Operations: Range Search?

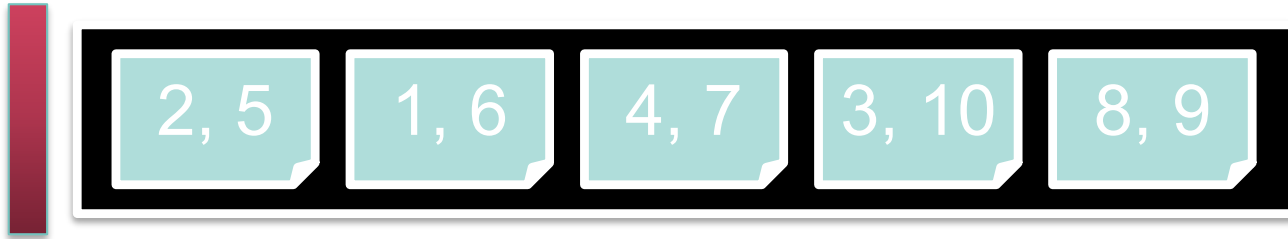
	Heap File	Sorted File
Scan all records	$B * D$	$B * D$
Equality Search	$0.5 * B * D$	$(\log_2 B) * D$
Range Search		
Insert		
Delete		

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk block

(all the)

Find Keys Between 7 and 9: Heap File

Heap File



- Always touch all blocks. Why?

Don't know
7 (How many) 9

Find Keys Between 7 and 9: Comparison

Heap File



- Find beginning of range

*binary to
find start*

Sorted File



- Search for start of range
- Scan right

Cost of Operations: Range Search Cost

	Heap File	Sorted File
Scan all records	$B * D$	$B * D$
Equality Search	$0.5 * B * D$	$(\log_2 B) * D$
Range Search	$B * D$	$\frac{((\log_2 B) + \text{pages}) * D}{D}$
Insert		
Delete		

→ # of pages matched

- **B**: The number of data blocks
- **R**: Number of records per block
- **D**: Average time to read/write disk block

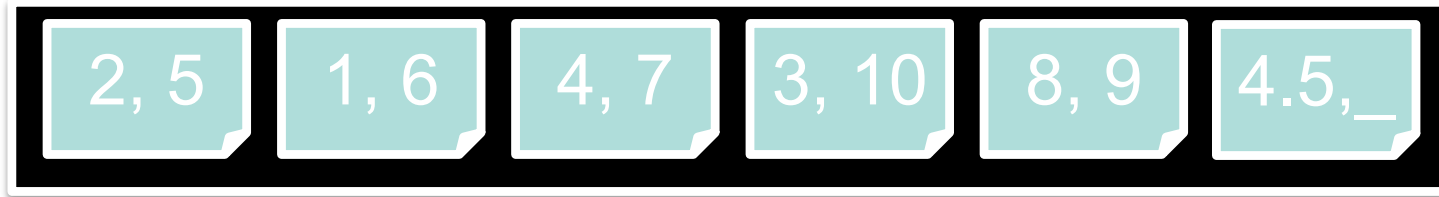
Cost of Operations: Insert?

	Heap File	Sorted File
Scan all records	$B * D$	$B * D$
Equality Search	$0.5 * B * D$	$(\log_2 B) * D$
Range Search	$B * D$	$((\log_2 B) + \text{pages}) * D$
Insert		
Delete		

- **B**: The number of data blocks
- **R**: Number of records per block
- **D**: Average time to read/write disk block

Insert 4.5: Heap File

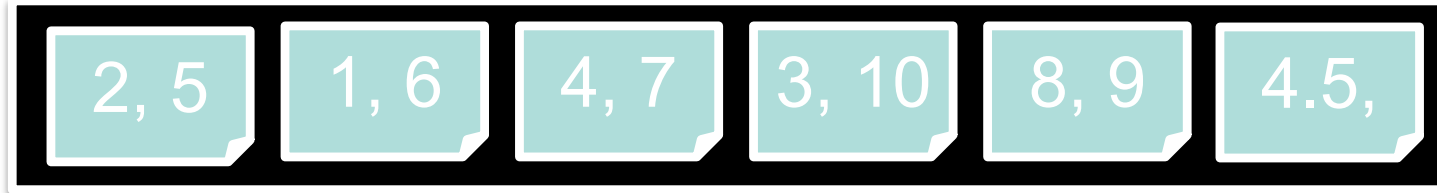
Heap File



- Stick at end of file
- Cost = $2 \times D$
- Why 2? double cost: read & write

Insert 4.5: Heap VS Sorted File

Heap File



- Read last page, append, write. Cost = $2 \cdot D$

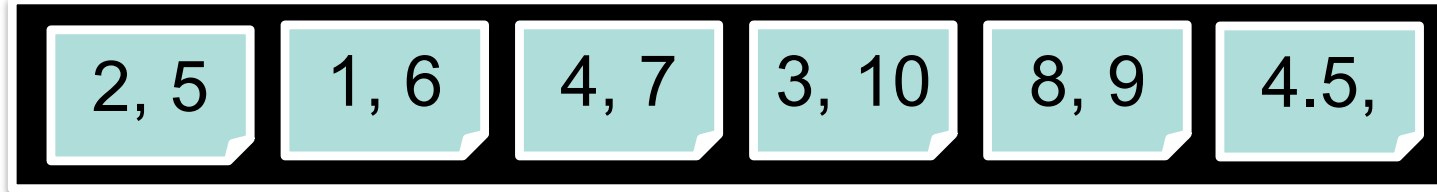
Sorted File



- Find location for record. Cost = $\log_2 BD$

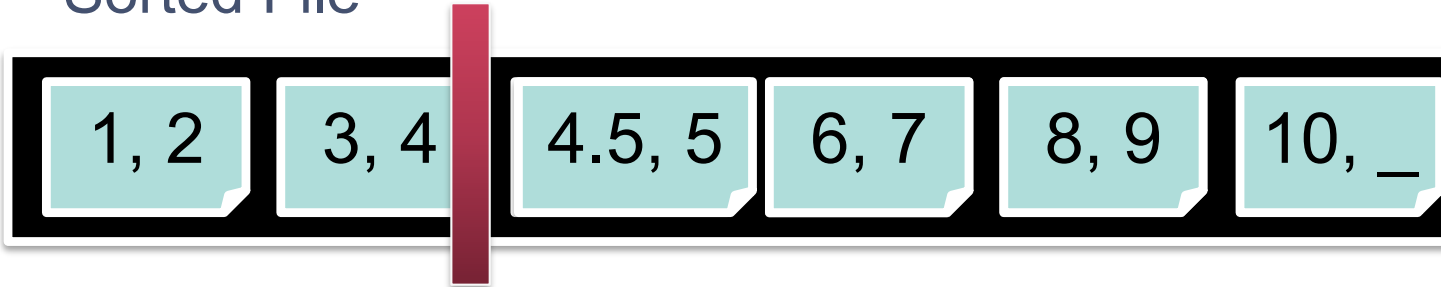
Insert 4.5: Heap Vs Sorted Pt 2

Heap File



- Read last page, append, write. Cost = $2 \cdot D$

Sorted File



- Find location for record. Cost = $\log_2 BD$
- Insert and shift rest of file *read and write*

Cost of Operations: Insert Cost

	Heap File	Sorted File
Scan all records	$B \cdot D$	$B \cdot D$
Equality Search	$0.5 \cdot B \cdot D$	$(\log_2 B) \cdot D$
Range Search	$B \cdot D$	$((\log_2 B) + \text{pages}) \cdot D$
Insert	$2 \cdot D$	$((\log_2 B) + B) \cdot D$
Delete		

$D \log_2 B + B D$

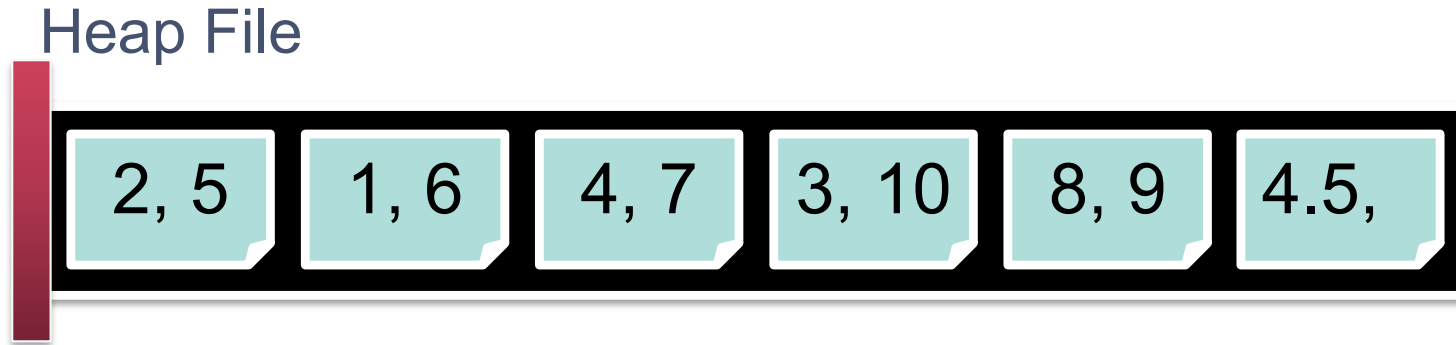
- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk block

Cost of Operations: Delete?

	Heap File	Sorted File
Scan all records	$B * D$	$B * D$
Equality Search	$0.5 * B * D$	$(\log_2 B) * D$
Range Search	$B * D$	$((\log_2 B) + \text{pages}) * D$
Insert	$2 * D$	$((\log_2 B) + B) * D$
Delete		

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk block

Delete 4.5: Heap File

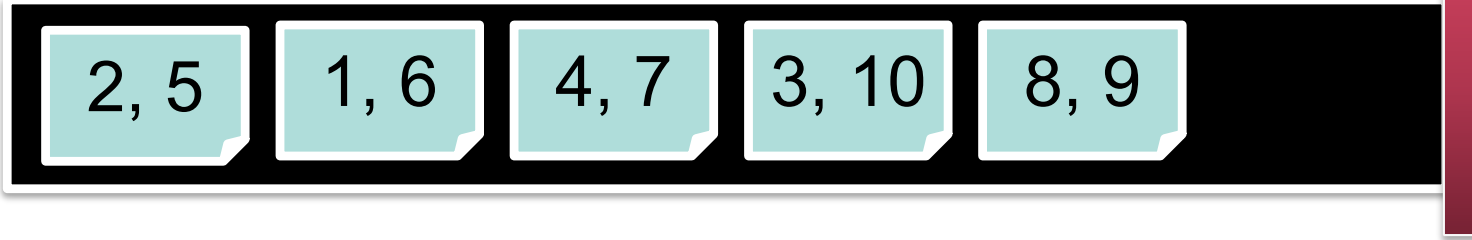


- Average case to find the record: **$B/2$ reads**
- Delete record from page
- Cost = $(B/2 + 1) * D$
 - Why + 1?

→ write of the modified.

Delete 4.5: Heap File Vs Sorted File

Heap File



- Average case runtime: $(B/2+1) * D$

Sorted File

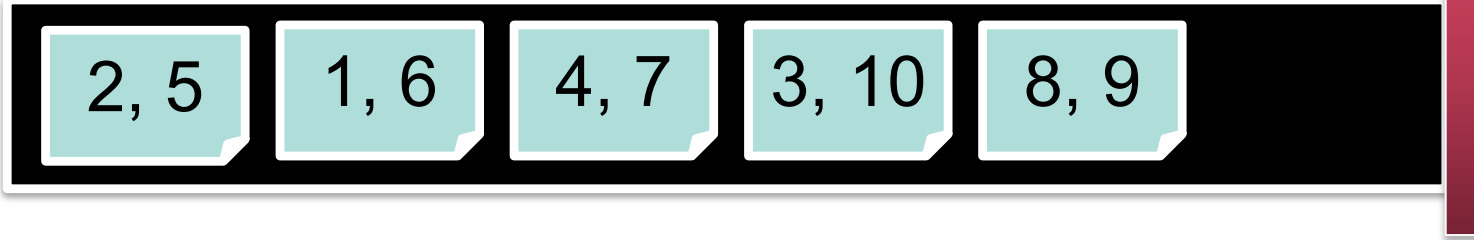


- Find location for record: $\log_2 B$
- Delete record in page \rightarrow Gap

(whole) but packed!
shift

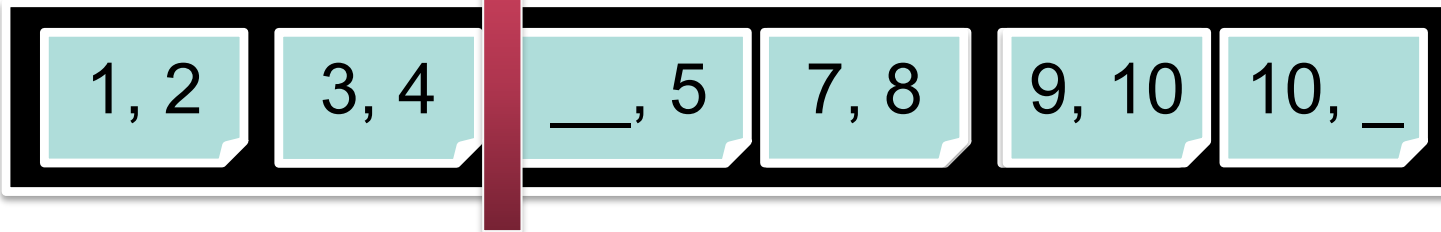
Delete 4.5: Heap File Vs Sorted File Pt 2

Heap File



- Average case runtime: $(B/2+1) * D$

Sorted File



- Find location for record: $\log_2 B$
- Shift the rest by 1 record $(2) * (B/2)$

read, write.

Cost of Operations Complete

	Heap File	Sorted File
Scan all records	$B * D$	$B * D$
Equality Search	$0.5 * B * D$	$(\log_2 B) * D$
Range Search	$B * D$	$\frac{((\log_2 B) + \text{pages}) * D}{D}$
Insert	$2 * D$	$((\log_2 B) + B) * D$
Delete	$(0.5 * B + 1) * D$	$((\log_2 B) + B) * D$

$\log_2 B + B$

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk block

Cost of Operations Complete Pt 2

	Heap File	Sorted File
Scan all records	$B * D$	$B * D$
Equality Search	$0.5 * B * D$	$(\log_2 B) * D$
Range Search	$B * D$	$((\log_2 B) + \text{pages}) * D$
Insert	$2 * D$	$((\log_2 B) + B) * D$
Delete	$(0.5 * B + 1) * D$	$((\log_2 B) + B) * D$

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk block
- Can we do better?
 - Indexes!

