

CS150A Database

Lu Sun

School of Information Science and Technology

ShanghaiTech University

Oct. 11, 2022

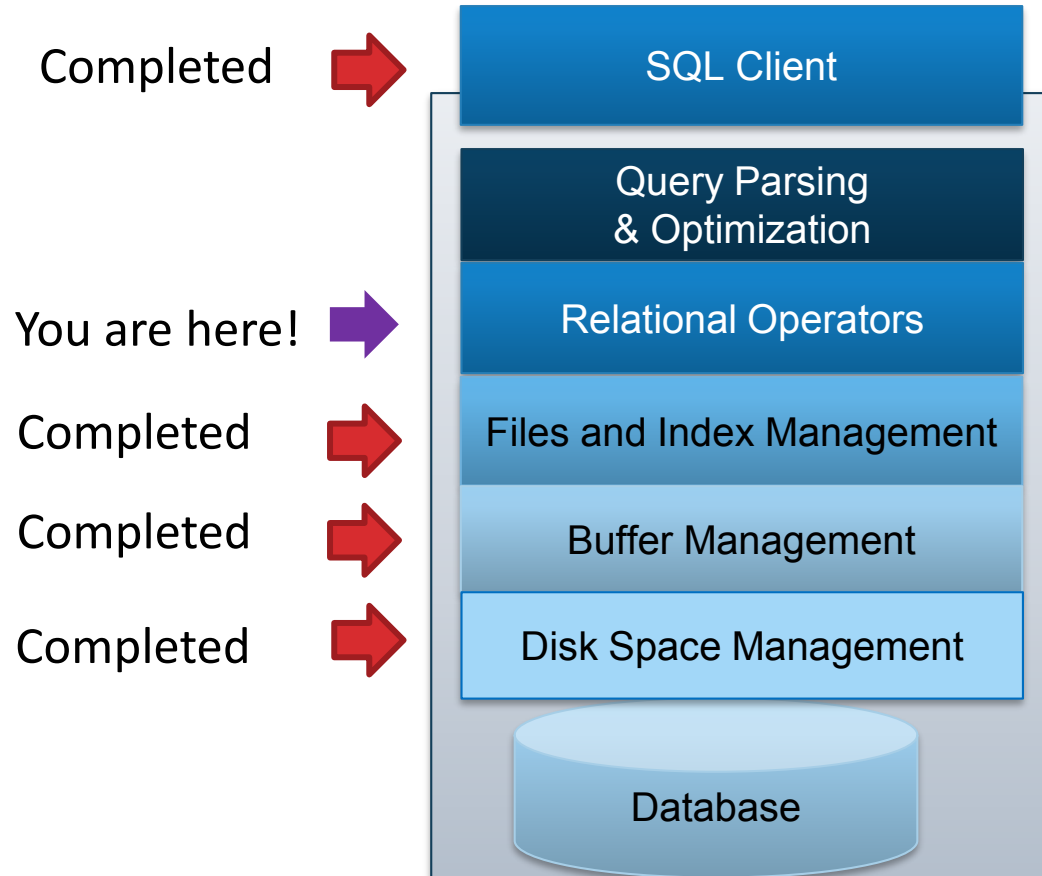
Today:

- Relational Algebra:
 - Basic Operators
 - Compound Operators

Readings:

- Database Management Systems (DBMS), Chapters 4.1-4.2

Architecture of a DBMS: What we've learned



Today: *definitions* of the relational operators.

Coming soon:
implementations

An Overview of the Layer Above

SQL Query

```
SELECT S.name
FROM Reserves R, Sailors S
WHERE R.sid = S.sid
AND R.bid = 100
AND S.rating > 5
```

Query Parser
& Optimizer

Relational Algebra

$$\pi_{S.name}(\sigma_{bid=100 \wedge rating > 5}(\text{Reserves} \bowtie_{R.sid=S.sid} \text{Sailors}))$$

Equivalent to...

Optimized (Physical) Query Plan:

On-the-fly
Project Iterator

On-the-fly
Select Iterator


Indexed Nested
Loop Join Iterator

Heap Scan
Iterator

But actually will
produce...


Operator Code

B+-Tree
Indexed Scan
Iterator



Tree of opt

(Logical) Query Plan:



SQL vs Relational Algebra

SQL Query

```
SELECT S.name
FROM Reserves R, Sailors S
WHERE R.sid = S.sid
AND R.bid = 100
AND S.rating > 5
```

Query Parser
& Optimizer

Relational Algebra

$$\pi_{S.name}(\sigma_{bid=100 \wedge rating > 5}(\text{Reserves} \bowtie_{R.sid=S.sid} \text{Sailors}))$$

Equivalent to...


SQL

A **declarative** expression
of the query result

But actually will
produce...

Relational Algebra

Operational description of
a computation.



Operator Code

B+-Tree
Indexed
Iterator

Systems execute relational algebra
query plan.

exhaustive

SQL (Structured Query Language)

```
SELECT S.name  
FROM Reserves R, Sailors S  
WHERE R.sid = S.sid  
AND R.bid = 100  
AND S.rating > 5
```

- Key **System** Features: *Why do we like SQL*
 - Declarative:
 - Say *what* you want, not *how* to get it
 - Enables system to optimize the *how*
- Foundation in formal Query Languages
 - **Relational Calculus**

History: Formal Relational QL's

- **Relational Calculus:** (Basis for SQL)

- Describe the result of computation
- Based on first order logic
- Tuple Relational Calculus (**TRC**)
 - $\{S \mid S \in \text{Sailors} \exists R \in \text{Reserves} (R.\text{sid} = S.\text{sid} \wedge R.\text{bid} = 103)\}$



Are these equivalent?


Can we go from one to the other?

- **Relational Algebra:**

- Algebra on sets
- Operational description of transformations

Codd's Theorem

- Established equivalence in expressivity between :
 - Relational Calculus
 - Relational Algebra
- Why an important result?
 - Connects declarative representation of queries with operational description
 - Constructive: we can compile SQL into relational algebra



Edgar F. "Ted" Codd
(1923 - 2003)
Turing Award 1981

declarative \longrightarrow operational

Relational Algebra Preliminaries



- Algebra of operators on relation instances
- $\pi_{S.name}(\sigma_{R.bid=100 \wedge S.rating>5}(R \bowtie_{R.sid=S.sid} S))$
 - Closed: result is also a relation instance
 - Enables rich composition!
 - Typed: input schema determines output
 - Can statically check whether queries are legal.

Relational Algebra and Sets

- Pure relational algebra has set semantics
 - No duplicate tuples in a relation instance
 - vs. SQL, which has multiset (bag) semantics
 - We will switch to multiset in the system discussion

no-dup sets

Relational Algebra Operators: Unary

- Unary Operators: on **single relation**
- **Projection** (π): Retains only desired columns (vertical) 
- **Selection** (σ): Selects a subset of rows (horizontal) 
- **Renaming** (ρ): Rename attributes and relations.

Relational Algebra Operators: Binary

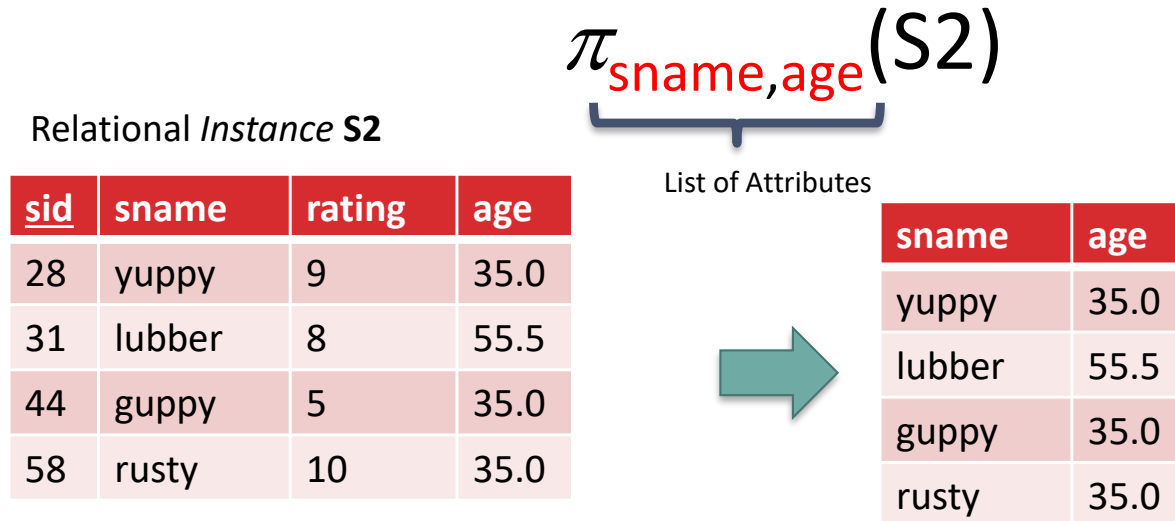
- Binary Operators: on **pairs of relations**
- **Union** (\cup): Tuples in r_1 or in r_2 .
- **Set-difference** ($-$): Tuples in r_1 , but not in r_2 .
- **Cross-product** (\times): Allows us to combine two relations.

Relational Algebra Operators: Compound

- Compound Operators: common “*macros*” for the above
- **Intersection** (\cap): **Tuples** in r_1 and in r_2 .
- **Joins** (\bowtie_{θ} , \bowtie): **Combine** relations that satisfy predicates

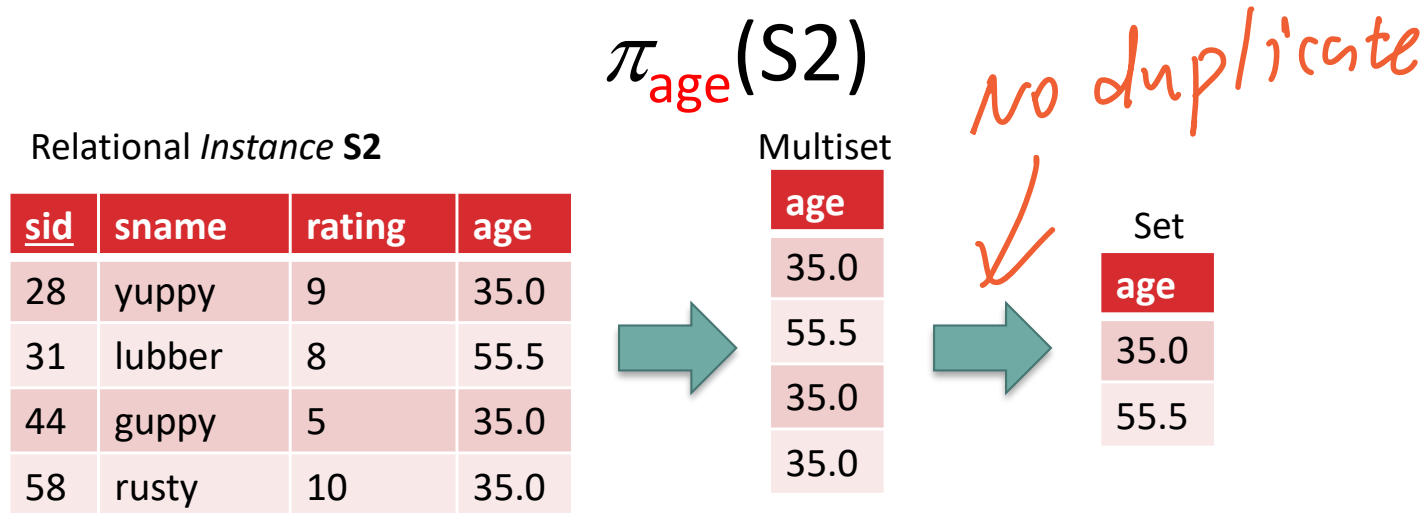
Projection (π) *vertical select*

- Corresponds to the SELECT list in SQL
- Schema determined by schema of attribute list *(integer, text...)*
 - Names and types correspond to input attributes
- *Selects a subset of columns (vertical)*



Projection (π), cont.

- Set semantics \rightarrow results in fewer rows
 - Real systems don't automatically remove duplicates
 - Why? (Semantics and Performance reasons)
 - *# maybe useful* *don't care.* Selects a subset of columns (vertical)



Selection(σ) *row*


- Corresponds to the WHERE clause in SQL
- Output schema same as input
- Duplicate Elimination? **Not needed.**
- *Selects a subset of rows (horizontal)*

$\sigma_{\text{rating} > 8}(\text{S2})$

Relational Instance S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

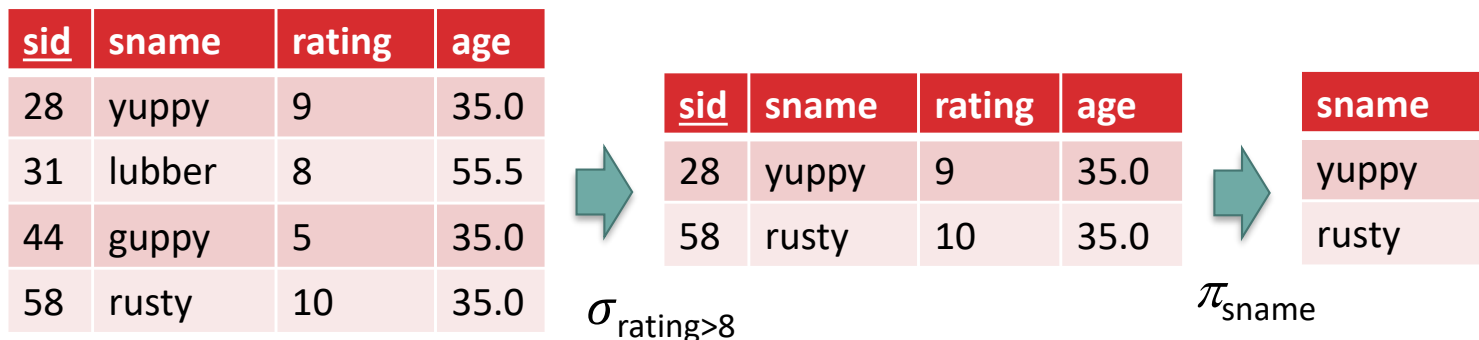
Selection Condition (Boolean Expression)



<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

Composing Select and Project

- Names of sailors with rating > 8: $\pi_{\text{sname}}(\sigma_{\text{rating}>8}(S2))$




- What about: $\sigma_{\text{rating}>8}(\pi_{\text{sname}}(S2))$
 - Invalid types. Input to $\sigma_{\text{rating}>8}$ does not contain rating.
- can not commute
不可交换

Union (U)

- Two input relations, must be *compatible*:
 - Same number of fields
 - Fields in corresponding positions have same type
- SQL Expression: UNION

$s1 \cup s2$



Union (U) VS Union ALL

- Duplicate elimination in practice?
- SQL's UNION vs UNION ALL



S1 U S2

<u>sid</u>	sname	rating	age
22	dustin	7	45
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Relational Instance S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

plus S2 to S1


Relational Instance S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

rm dup

Set Difference (-)

- Same as with union, both input relations must be *compatible*.
- SQL Expression: EXCEPT



Set Difference (-), cont.

- Duplicate elimination?
 - Not required
- EXCEPT vs EXCEPT ALL

Relational Instance S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Relational Instance S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S1 - S2

<u>sid</u>	sname	rating	age
22	dustin	7	45

S2 - S1

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
44	guppy	5	35.0

Cross-Product (\times)

- $R1 \times S1$: Each row of $R1$ paired with each row of $S1$

R1:			S1:				$R_1 \times S_1$								
<u>sid</u>	<u>bid</u>	<u>day</u>	×	<u>sid</u>	sname	rating	age	=	sid	bid	day	sid	sname	rating	age
22	101	10/10/96		22	dustin	7	45.0		22	101	10/10/96	22	dustin	7	45.0
58	103	11/12/96		31	lubber	8	55.5		22	101	10/10/96	31	lubber	8	55.5
				58	rusty	10	35.0		58	103	11/12/96	58	rusty	10	35.0
									22	101	10/10/96	58	rusty	10	35.0
									58	103	11/12/96	22	dustin	7	45.0
									58	103	11/12/96	31	lubber	8	55.5
									58	103	11/12/96	58	rusty	10	35.0

- How many rows in result? $|R1| \times |R2|$
- Schema compatibility? Not needed.
- Duplicates? None generated.

Renaming (ρ = “rho”)

- *Renames relations and their attributes:*
- Note that relational algebra doesn't require names.
 - We could just use positional arguments.

pos name

$$\rho(\text{Temp1}(1 \rightarrow \text{sid1}, 4 \rightarrow \text{sid2}), R1 \times S1)$$

Output Relation
Name

Renaming List
position \rightarrow New Name

Input
Relation

$R1 \times S1$

sid	bid	day	sid	sname	rating	age
22	101	10/10/96	22	dustin	7	45.0
22	101	10/10/96	31	lubber	8	55.5
22	101	10/10/96	58	rusty	10	35.0
58	103	11/12/96	22	dustin	7	45.0
58	103	11/12/96	31	lubber	8	55.5
58	103	11/12/96	58	rusty	10	35.0




Temp1

sid1	bid	day	sid2	sname	rating	age
22	101	10/10/96	22	dustin	7	45.0
22	101	10/10/96	31	lubber	8	55.5
22	101	10/10/96	58	rusty	10	35.0
58	103	11/12/96	22	dustin	7	45.0
58	103	11/12/96	31	lubber	8	55.5
58	103	11/12/96	58	rusty	10	35.0

Compound Operator: Intersection

- Same as with union, both input relations must be *compatible*.
- SQL Expression: INTERSECT

$$S1 \cap S2$$



Intersection (\cap)

- Equivalent to:
 $S1 - (S1 - S2)$



Relational Instance **S1**

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Relational Instance **S2**


<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$S1 \cap S2$

<u>sid</u>	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0



Intersection (\cap), Pt 2

- $S1 \cap S2 = ?$





Intersection (\cap), Pt 3

- $S1 \cap S2 = S1 - ?$



Intersection (\cap), Pt 4


- $S1 \cap S2 = S1 - (S1 - S2)$



- Is intersection monotonic?
单调 *"only"*

$$R_1 \subseteq R_2 \Rightarrow S \cap R_1 \subseteq S \cap R_2$$

Compound Operator: Join

- Joins are compound operators (like intersection):
 - Generally, $\sigma_{\theta}(R \times S)$
 - **Hierarchy** of common kinds:
 - **Theta Join** (\bowtie_{θ}): join on logical expression θ
 - **Equi-Join**: theta join with theta being a conjunction of equalities
 - **Natural Join** (\bowtie): equi-join on all matching column names
- 

Note: we will need to learn a good join algorithm.

Avoid cross-product if we can!!

Theta Join (\bowtie_{θ}) Example

- $R1 \bowtie_{\text{sid=sid}} S1$
logical expression

R1:

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

S1:

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

$\bowtie_{\text{sid=sid}}$

=

sid	bid	day	sid	sname	rating	age
22	101	10/10/96	22	dustin	7	45.0
58	103	11/12/96	58	rusty	10	35.0

- Note that output needs a **rename** operator!

by pos

Another Theta Join (\bowtie_{θ}) Example do condition on cross-product outcome

- $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$
- **Example:** *More senior sailors for each sailor.*
- $S1 \bowtie_{f4 < f8} S1$

f1	f2	f3	f4
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1				S1			
f1	f2	f3	f4	f5	f6	f7	f8
22	dustin	7	45.0	22	dustin	7	45.0
22	dustin	7	45.0	31	lubber	8	55.5
22	dustin	7	45.0	58	rusty	10	35.0
31	lubber	8	55.5	22	dustin	7	45.0
31	lubber	8	55.5	31	lubber	8	55.5
31	lubber	8	55.5	58	rusty	10	35.0
58	rusty	10	35.0	22	dustin	7	45.0
58	rusty	10	35.0	31	lubber	8	55.5
58	rusty	10	35.0	58	rusty	10	35.0

→ cross product outcome

Another Theta Join (\bowtie_{θ}), Pt 2

- $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$
- **Example:** *More senior sailors for each sailor.*
- $S1 \bowtie_{\text{age} < \text{age2}} S1$

S1:

f1	f2	f3	f4
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

(f4 < f8) on result

S1				S1			
f1	f2	f3	f4	f5	f6	f7	f8
22	dustin	7	45.0	22	dustin	7	45.0
22	dustin	7	45.0	31	lubber	8	55.5
22	dustin	7	45.0	58	rusty	10	35.0
31	lubber	8	55.5	22	dustin	7	45.0
31	lubber	8	55.5	31	lubber	8	55.5
31	lubber	8	55.5	58	rusty	10	35.0
58	rusty	10	35.0	22	dustin	7	45.0
58	rusty	10	35.0	31	lubber	8	55.5
58	rusty	10	35.0	58	rusty	10	35.0

Another Theta Join (\bowtie_{θ}), Pt 3

- $\mathbf{R} \bowtie_{\theta} \mathbf{S} = \sigma_{\theta}(\mathbf{R} \times \mathbf{S})$
- **Example:** *More senior sailors for each sailor.*
- $\mathbf{S1} \bowtie_{f4 < f8} \mathbf{S1}$

S1:

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1				S1			
sid	sname	rating	age	sid	sname	rating	age2
22	dustin	7	45.0	31	lubber	8	55.5
58	rusty	10	35.0	22	dustin	7	45.0
58	rusty	10	35.0	31	lubber	8	55.5

- Result *schema* same as that of cross-product.
- Special Case:
 - **Equi-Join:** *theta join with AND of = predicates*
 - Special special case **Natural Join** ...

Natural Join (\bowtie)

- Special case of **equi-join** in which equalities are specified for all matching fields and duplicate fields are projected away

$$\mathbf{R} \bowtie \mathbf{S} = \pi_{\text{unique fld.}} \sigma_{\text{eq. matching fld.}} (\mathbf{R} \times \mathbf{S})$$

- Compute $\mathbf{R} \times \mathbf{S}$
- Select rows where **fields appearing in both relations** have equal values
- Project onto the set of all unique fields.

Natural Join (\bowtie) Pt 2

- $R \bowtie S = \pi_{\text{unique fld.}} \sigma_{\text{eq. matching fld.}} (R \times S)$

Semantic

*matching: name
and value eq*

$R1 \bowtie S1$

sid	bid	day	sid	sname	rating	age
22	101	10/10/96	22	dustin	7	45.0
22	101	10/10/96	31	lubber	8	55.5
22	101	10/10/96	58	rusty	10	35.0
58	103	11/12/96	22	dustin	7	45.0
58	103	11/12/96	31	lubber	8	55.5
58	103	11/12/96	58	rusty	10	35.0

R1:

sid	bid	day
22	101	10/10/96
58	103	11/12/96

S1:

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Natural Join (\bowtie), Pt 3

- ↓ →
- $R \bowtie S = \pi_{\text{unique fld.}} \sigma_{\text{eq. matching fld.}} (R \times S)$

*remove
dup matching*

$R1 \bowtie S1$

sid	bid	day	sname	rating	age
22	101	10/10/96	dustin	7	45.0
58	103	11/12/96	rusty	10	35.0

R1:

sid	bid	day
22	101	10/10/96
58	103	11/12/96

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

- Commonly used for foreign key joins (as above).

Extended Relational Algebra

- **Group By / Aggregation Operator (γ):**
 - $\gamma_{\text{age, AVG(rating)}}(\text{Sailors})$
 - With selection (HAVING clause):
 - $\gamma_{\text{age, AVG(rating), COUNT(*) > 2}}(\text{Sailors})$
 - Textbook uses two operators:
 - GROUP BY age, AVG(rating) (Sailors)
 - HAVING COUNT(*) > 2
(GROUP BY age, AVG(rating)(Sailors))
- Handwritten notes:*
An orange arrow points from the γ in the first bullet to the γ in the second bullet.
The symbol $\sigma \rho \tau$ is written in orange above the second bullet.
The symbol $\sigma \rho \tau$ is written in orange below the second bullet.

Summary

- Relational Algebra: a small set of operators mapping relations to relations
 - Operational, in the sense that you specify the explicit order of operations
 - A closed set of operators! Mix and match.
- Basic ops include: σ , π , \times , \cup , $-$
- Important compound ops: \cap , \bowtie